

Predicting Loan Repayment Using Machine Learning: **Decision Tree vs. Random Forest**

Exploring publicly available data from [LendingClub.com](https://lendingclub.com). Lending data from 2007-2010 and trying to classify and predict whether or not the borrower paid back their loan in full.

Here are what the columns represent:

- credit.policy: 1 if the customer meets the credit underwriting criteria of LendingClub.com, and 0 otherwise.
- purpose: The purpose of the loan (takes values "credit_card", "debt_consolidation", "educational", "major_purchase", "small_business", and "all_other").
- int.rate: The interest rate of the loan, as a proportion (a rate of 11% would be stored as 0.11). Borrowers judged by LendingClub.com to be more risky are assigned higher interest rates.
- installment: The monthly installments owed by the borrower if the loan is funded.
- log.annual.inc: The natural log of the self-reported annual income of the borrower.
- dti: The debt-to-income ratio of the borrower (amount of debt divided by annual income).
- fico: The FICO credit score of the borrower.
- days.with.cr.line: The number of days the borrower has had a credit line.
- revol.bal: The borrower's revolving balance (amount unpaid at the end of the credit card billing cycle).
- revol.util: The borrower's revolving line utilization rate (the amount of the credit line used relative to total credit available).
- inq.last.6mths: The borrower's number of inquiries by creditors in the last 6 months.
- delinq.2yrs: The number of times the borrower had been 30+ days past due on a payment in the past 2 years.
- pub.rec: The borrower's number of derogatory public records (bankruptcy filings, tax liens, or judgments).

DATA SETUP

- Importing numpy, pandas, visualisation libraries.

[1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

- Reading csv file

[5]:

```
loans=pd.read_csv('loan_data.csv')
```

- Information of the file

[7]:

```
loans.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   credit.policy          9578 non-null   int64
1   purpose                9578 non-null   object
2   int.rate               9578 non-null   float64
3   installment            9578 non-null   float64
4   log.annual.inc         9578 non-null   float64
5   dti                    9578 non-null   float64
6   fico                   9578 non-null   int64
7   days.with.cr.line      9578 non-null   float64
8   revol.bal              9578 non-null   int64
9   revol.util             9578 non-null   float64
10  inq.last.6mths         9578 non-null   int64
11  delinq.2yrs            9578 non-null   int64
12  pub.rec                9578 non-null   int64
13  not.fully.paid         9578 non-null   int64
dtypes: float64(6), int64(7), object(1)
memory usage: 1.0+ MB
```

[9]: `loans.describe()`

	credit.policy	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	delinq.2yrs	pub.rec	not.fully.paid
count	9578.000000	9578.000000	9578.000000	9578.000000	9578.000000	9578.000000	9578.000000	9.578000e+03	9578.000000	9578.000000	9578.000000	9578.000000	9578.000000
mean	0.804970	0.122640	319.089413	10.932117	12.606679	710.846314	4560.767197	1.691396e+04	46.799236	1.577469	0.163708	0.063708	0.063708
std	0.396245	0.026847	207.071301	0.614813	6.883970	37.970537	2496.930377	3.375619e+04	29.014417	2.200245	0.546215	0.263708	0.263708
min	0.000000	0.060000	15.670000	7.547502	0.000000	612.000000	178.958333	0.000000e+00	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	0.103900	163.770000	10.558414	7.212500	682.000000	2820.000000	3.187000e+03	22.600000	0.000000	0.000000	0.000000	0.000000
50%	1.000000	0.122100	268.950000	10.928884	12.665000	707.000000	4139.958333	8.596000e+03	46.300000	1.000000	0.000000	0.000000	0.000000
75%	1.000000	0.140700	432.762500	11.291293	17.950000	737.000000	5730.000000	1.824950e+04	70.900000	2.000000	0.000000	0.000000	0.000000
max	1.000000	0.216400	940.140000	14.528354	29.960000	827.000000	17639.958330	1.207359e+06	119.000000	33.000000	13.000000	5.000000	5.000000

- Sample dataset

[11]: `loans.head()`

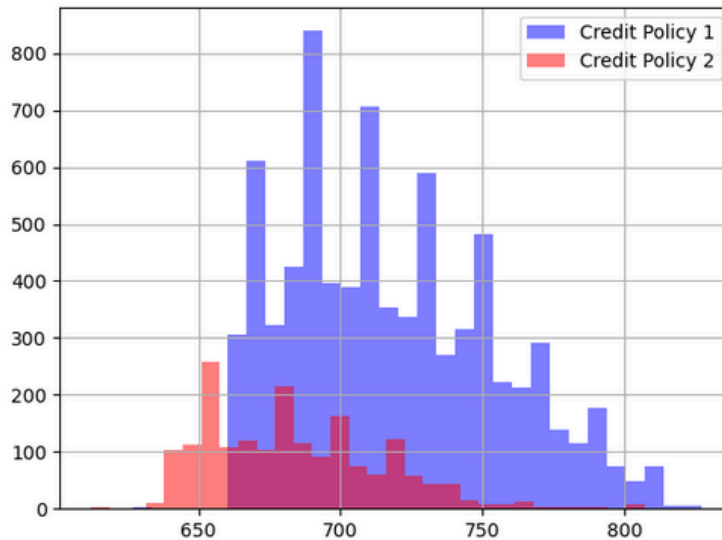
	credit.policy	purpose	int.rate	installment	log.annual.inc	dti	fico	days.with.cr.line	revol.bal	revol.util	inq.last.6mths	delinq.2yrs	pub.rec	not.fully.paid
0	1	debt_consolidation	0.1189	829.10	11.350407	19.48	737	5639.958333	28854	52.1	0	0	0	0
1	1	credit_card	0.1071	228.22	11.082143	14.29	707	2760.000000	33623	76.7	0	0	0	0
2	1	debt_consolidation	0.1357	366.86	10.373491	11.63	682	4710.000000	3511	25.6	1	0	0	0
3	1	debt_consolidation	0.1008	162.34	11.350407	8.10	712	2699.958333	33667	73.2	1	0	0	0
4	1	credit_card	0.1426	102.92	11.299732	14.97	667	4066.000000	4740	39.5	0	1	0	0

EXPLORATORY DATA ANALYSIS

- Creating a histogram of two FICO distributions on top of each other, one for each credit.policy outcome and not.fully.paid column.

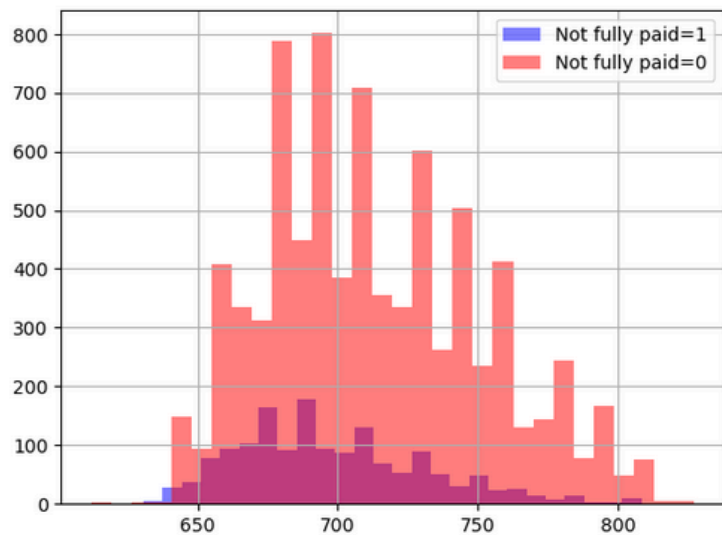
```
[41]: loans[loans['credit.policy']==1]['fico'].hist(alpha=0.5,color='blue',bins=30,label='Credit Policy 1')
      loans[loans['credit.policy']==0]['fico'].hist(alpha=0.5,color='red',bins=30,label='Credit Policy 2')
      plt.legend()
```

```
[41]: <matplotlib.legend.Legend at 0x1ef28522ed0>
```



```
[47]: loans[loans['not.fully.paid']==1]['fico'].hist(alpha=0.5,bins=30,color='blue',label='Not fully paid=1')
      loans[loans['not.fully.paid']==0]['fico'].hist(alpha=0.5,bins=30,color='red',label='Not fully paid=0')
      plt.legend()
```


```
[47]: <matplotlib.legend.Legend at 0x1ef2cde2ed0>
```



- Creating a countplot using seaborn showing the counts of loans by purpose, with the colour hue defined by not.fully.paid.

```
[ ]: plt.figure(figsize=(11,7))
sns.countplot(x='purpose',hue='not.fully.paid',data=loans,palette='Set1')
```


```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x119996828>
```



- The trend between FICO score and interest rate.

```
[55]: sns.jointplot(data=loans,x='fico',y='int.rate',color='purple')
```


```
[55]: <seaborn.axisgrid.JointGrid at 0x1ef2d0b1750>
```



- Creating a Implots to see if the trend differed between not.fully.paid and credit.policy.

```
[63]: sns.lmplot(data=loans,x='fico',y='int.rate',hue='credit.policy',col='not.fully.paid',palette='Set1')
```

```
[63]: <seaborn.axisgrid.FacetGrid at 0x1ef2c982110>
```



Categorical Features

- **Purpose** column is categorical, transforming them using dummy variables so sklearn will be able to understand them.

```
63]: cat_feats=['purpose']
final_data = pd.get_dummies(loans,columns=cat_feats,drop_first=True)
final_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9578 entries, 0 to 9577
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                ---
0   credit.policy                         9578 non-null   int64
1   int.rate                             9578 non-null   float64
2   installment                          9578 non-null   float64
3   log.annual.inc                       9578 non-null   float64
4   dti                                  9578 non-null   float64
5   fico                                 9578 non-null   int64
6   days.with.cr.line                    9578 non-null   float64
7   revol.bal                            9578 non-null   int64
8   revol.util                           9578 non-null   float64
9   inq.last.6mths                       9578 non-null   int64
10  delinq.2yrs                          9578 non-null   int64
11  pub.rec                              9578 non-null   int64
12  not.fully.paid                       9578 non-null   int64
13  purpose_credit_card                  9578 non-null   bool
14  purpose_debt_consolidation           9578 non-null   bool
15  purpose_educational                  9578 non-null   bool
16  purpose_home_improvement             9578 non-null   bool
17  purpose_major_purchase               9578 non-null   bool
18  purpose_small_business                9578 non-null   bool
dtypes: bool(6), float64(6), int64(7)
memory usage: 1.0 MB
```

In the info ... purpose has categorical data now

Train Test Split

- Using sklearn to split the data into a training set and a testing set.

```
[107]: from sklearn.model_selection import train_test_split
```

```
[119]: X=final_data.drop('not.fully.paid',axis=1)
y=final_data['not.fully.paid']
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3)
```

Training a Decision Tree Model

- Importing DecisionTreeClassifier, creating an instance and fit it to the training data.

```
[165]: from sklearn.tree import DecisionTreeClassifier
dtree=DecisionTreeClassifier()
dtree.fit(X_train,y_train)
```

```
[165]: DecisionTreeClassifier
```

Prediction and Evaluation of Decision Tree

- Creating predictions from the test set and creating a classification report and a confusion matrix.

```
[167]: prediction=dtree.predict(X_test)
```

```
[169]: from sklearn.metrics import classification_report,confusion_matrix
```

```
[171]: print(classification_report(y_test,prediction))
```

	precision	recall	f1-score	support
0	0.85	0.85	0.85	2405
1	0.24	0.24	0.24	469
accuracy			0.75	2874
macro avg	0.54	0.54	0.54	2874
weighted avg	0.75	0.75	0.75	2874

```
[173]: print(confusion_matrix(y_test,prediction))
```

```
[[2038 367]
 [ 356 113]]
```

Training a Random Forest Model

- Importing DecisionTreeClassifier, creating an instance and fit it to the training data.

```
[175]: from sklearn.ensemble import RandomForestClassifier
```

```
[177]: rfc=RandomForestClassifier(n_estimators=200)
```

```
[179]: rfc.fit(X_train,y_train)
```

```
[179]: RandomForestClassifier
```

Prediction and Evaluation of Random Forest

- Creating predictions from the test set and creating a classification report and a confusion matrix.

```
[181]: pred_rfc=rfc.predict(X_test)
       print(classification_report(y_test,pred_rfc))
```

	precision	recall	f1-score	support
0	0.84	1.00	0.91	2405
1	0.53	0.02	0.04	469
accuracy			0.84	2874
macro avg	0.68	0.51	0.48	2874
weighted avg	0.79	0.84	0.77	2874

```
[183]: print(confusion_matrix(y_test,pred_rfc))
```

```
[[2396   9]
 [ 459  10]]
```

- **Preferred model:**

The Random Forest model, since it demonstrated superior performance compared to the Decision Tree model, achieving higher accuracy and better handling of overfitting. This improvement is attributed to the ensemble nature of the Random Forest, which leverages multiple decision trees to enhance predictive performance and robustness.